

Software Barrier Algorithms

INCOMPLETE
Shows Progress

Outline

- Hardware Background
- Software Barrier Algorithms Implemented
- Intel Sandy Bridge Architecture
- Experiments and Results
- Conclusions

Hardware Background (1/3)

- Correctness Related
 - Read/Write/Full fences required for correct execution
 - Used C/C++11 memory model to abstract the underlying hardware machine memory model
 - Portability: all compilers support the C/C++11 memory model (at least in their newest versions)
 - DETAILS FOR THE MEMORY MODEL OMITTED

Hardware Background (2/3)

- Performance Related

- Atomic Primitives Used

- Instead of locks I used the atomics.
 - `std::atomic<int> i`, declares an atomic integer variable `i`.
 - `i.load(memory_order)/i.store(v,memory_order)/i.exchange(v,memory_order)`
 - `i.compare_exchange_weak(vold,vnew,memory_order)/...`

- False-Sharing

- Even if threads use different data variables, the cache ping-pong effect can still arise if those variables lie inside the same cache-line.
 - HARDWARE PREFETCHING NOT YET DEALT WITH

- Relaxed Model

- The implementations use the Relaxed Ordering offered by the C/C++ memory model which map efficiently to the underlying memory model by providing fine-grained control over the amount of *ordering* and *synchronization* required around the memory accesses of different threads
 - What does this mean for performance?
 - See next slide

Hardware Background (3/3)

AArch64

| C/C++11 Operation | AArch64 implementation |
|-------------------|--|
| Load Relaxed: | LDR |
| Load Consume: | LDR + preserve dependencies until next kill_dependency OR LDAR |
| Load Acquire: | LDAR |
| Load Seq Cst: | LDAR |
| Store Relaxed: | STR |
| Store Release: | STLR |
| Store Seq Cst: | STLR |
| Cmpxchg Relaxed: | loop: ldxr roldval, [rptr]; cmp roldval, rold; bne _exit; stxr rres, rnewval, [rptr]; cbz rres, _loop; _exit |
| Cmpxchg Acquire: | loop: ldaxr roldval, [rptr]; cmp roldval, rold; beq _exit; stxr rres, rnewval, [rptr]; cbz rres, _loop; _exit |
| Cmpxchg Release: | loop: ldxr roldval, [rptr]; cmp roldval, rold; bne _exit; stlxr rres, rnewval, [rptr]; cbz rres, _loop; _exit |
| Cmpxchg AcqRel: | loop: ldaxr roldval, [rptr]; cmp roldval, rold; bne _exit; stlxr rres, rnewval, [rptr]; cbz rres, _loop; _exit |
| Cmpxchg SeqCst : | loop: ldaxr roldval, [rptr]; cmp roldval, rold; bne _exit; stlxr rres, rnewval, [rptr]; cbz rres, _loop; _exit |
| Acquire Fence: | DMB LD |
| Release Fence: | DMB |
| AcqRel Fence: | DMB |
| SeqCst Fence: | DMB |

IBM Power / PowerPC

| C/C++11 Operation | Power implementation |
|---------------------------|---|
| Load Relaxed: | ld |
| Load Consume: | ld + preserve dependencies until next kill_dependency OR ld; cmp; bc; isync |
| Load Acquire: | ld; cmp; bc; isync |
| Load Seq Cst: | hwsync; ld; cmp; bc; isync |
| Store Relaxed: | st |
| Store Release: | lwsync; st |
| Store Seq Cst: | hwsync; st |
| Cmpxchg Relaxed (32 bit): | loop: lwarx; cmp; bc _exit; stwcx.; bc _loop; _exit: |
| Cmpxchg Acquire (32 bit): | loop: lwarx; cmp; bc _exit; stwcx.; bc _loop; isync; _exit: |
| Cmpxchg Release (32 bit): | lwsync; _loop: lwarx; cmp; bc _exit; stwcx.; bc _loop; _exit: |
| Cmpxchg AcqRel (32 bit): | lwsync; _loop: lwarx; cmp; bc _exit; stwcx.; bc _loop; isync; _exit |
| Cmpxchg SeqCst (32 bit): | hwsync; _loop: lwarx; cmp; bc _exit; stwcx.; bc _loop; isync; _exit |
| Acquire Fence: | lwsync |
| Release Fence: | lwsync |
| AcqRel Fence: | lwsync |
| SeqCst Fence: | hwsync |

Itanium

| C/C++11 Operation | IA64 implementation |
|-------------------|---------------------|
| Load Relaxed: | ld.acq |
| Load Consume: | ld.acq |
| Load Acquire: | ld.acq |
| Load Seq Cst: | ld.acq |
| Store Relaxed: | st.rel |
| Store Release: | st.rel |
| Store Seq Cst: | st.rel; mf |
| Cmpxchg Acquire: | cmpxchg.acq |
| Cmpxchg Release: | cmpxchg.rel |
| Cmpxchg AcqRel: | cmpxchg.rel; mf |
| Cmpxchg SeqCst: | cmpxchg.rel; mf |
| Consume Fence: | <ignore> |
| Acquire Fence: | <ignore> |
| Release Fence: | <ignore> |
| Acq Rel Fence: | <ignore> |
| Seq Cst Fence: | mf |

| C/C++11 Operation | ARM implementation |
|---------------------------|---|
| Load Relaxed: | ldr |
| Load Consume: | ldr + preserve dependencies until next kill_dependency OR ldr; teq; beq; isb OR ldr; dmb |
| Load Acquire: | ldr; teq; beq; isb OR ldr; dmb |
| Load Seq Cst: | ldr; dmb |
| Store Relaxed: | str |
| Store Release: | dmb; str |
| Store Seq Cst: | dmb; str; dmb |
| Cmpxchg Relaxed (32 bit): | _loop: ldrex roldval, [rptr]; mov rres, 0; teq roldval, rold; strex rres, rnewval, [rptr]; teq rres, 0; bne _loop |
| Cmpxchg Acquire (32 bit): | _loop: ldrex roldval, [rptr]; mov rres, 0; teq roldval, rold; strex rres, rnewval, [rptr]; teq rres, 0; bne _loop; isb |
| Cmpxchg Release (32 bit): | dmb; _loop: ldrex roldval, [rptr]; mov rres, 0; teq roldval, rold; strex rres, rnewval, [rptr]; teq rres, 0; bne _loop; |
| Cmpxchg AcqRel (32 bit): | dmb; _loop: ldrex roldval, [rptr]; mov rres, 0; teq roldval, rold; strex rres, rnewval, [rptr]; teq rres, 0; bne _loop; isb |
| Cmpxchg SeqCst (32 bit): | dmb; _loop: ldrex roldval, [rptr]; mov rres, 0; teq roldval, rold; strex rres, rnewval, [rptr]; teq rres, 0; bne _loop; isb |
| Acquire Fence: | dmb |
| Release Fence: | dmb |
| AcqRel Fence: | dmb |
| SeqCst Fence: | dmb |

x86 (including x86-64)

| C/C++11 Operation | x86 implementation |
|-------------------|--|
| Load Relaxed: | MOV (from memory) |
| Load Consume: | MOV (from memory) |
| Load Acquire: | MOV (from memory) |
| Load Seq Cst: | MOV (from memory) |
| Store Relaxed: | MOV (into memory) |
| Store Release: | MOV (into memory) |
| Store Seq Cst: | (LOCK) XCHG // alternative: MOV (into memory),MFENCE |
| Consume Fence: | <ignore> |
| Acquire Fence: | <ignore> |
| Release Fence: | <ignore> |
| Acq Rel Fence: | <ignore> |
| Seq Cst Fence: | MFENCE |

Software Barrier Algorithms Implemented

- Centralized Sense Reversing Barrier
- Static Tree Barrier
 - Arrival Tree
 - Departure Tree
- Static Tree Barrier with Global Departure
 - Arrival Tree
 - Global atomic boolean flag for departure
 - Like the departure phase of the centralized sense-reversing barrier algorithm
- Dissemination Barrier
 - NOT YET

Intel Sandy Bridge Architecture

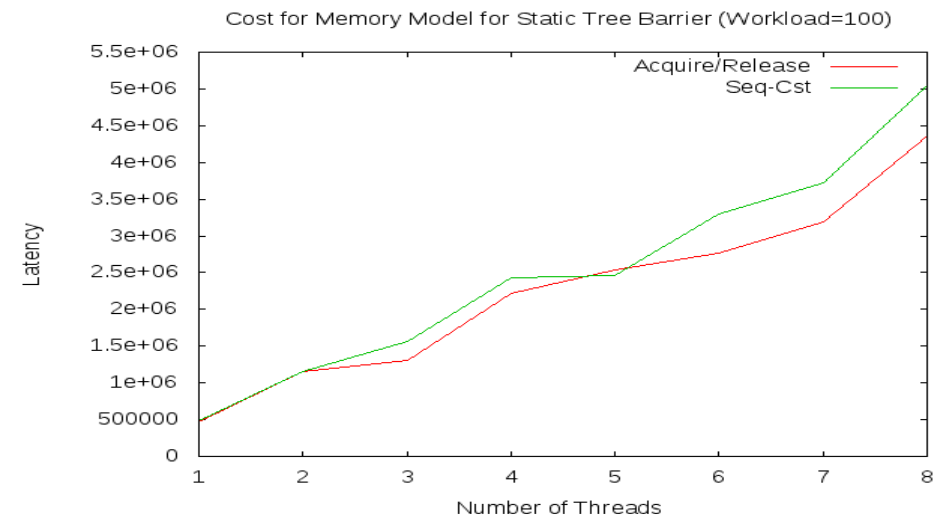
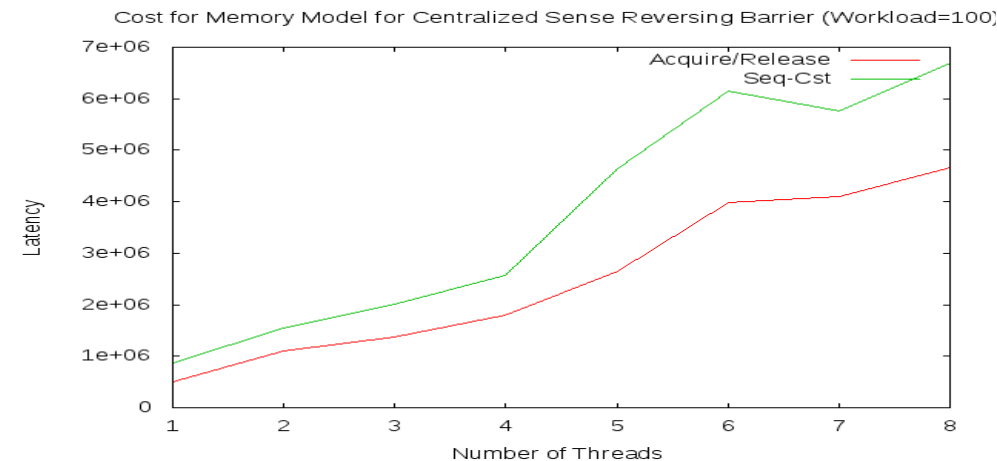
- OMITTED

Experiments and Results (1/5)

- Thread mapping:
 - First fill first thread inside each core from left to right
 - Then fill the second thread inside each core from left to right
- Experimental Process:
 - 10.000 barrier episodes. Each thread makes a random workload of 100 counter increments between episodes.
 - Repeat 30 times and measure (lower,mean,upper) *latency* with a confidence level of 99.9%

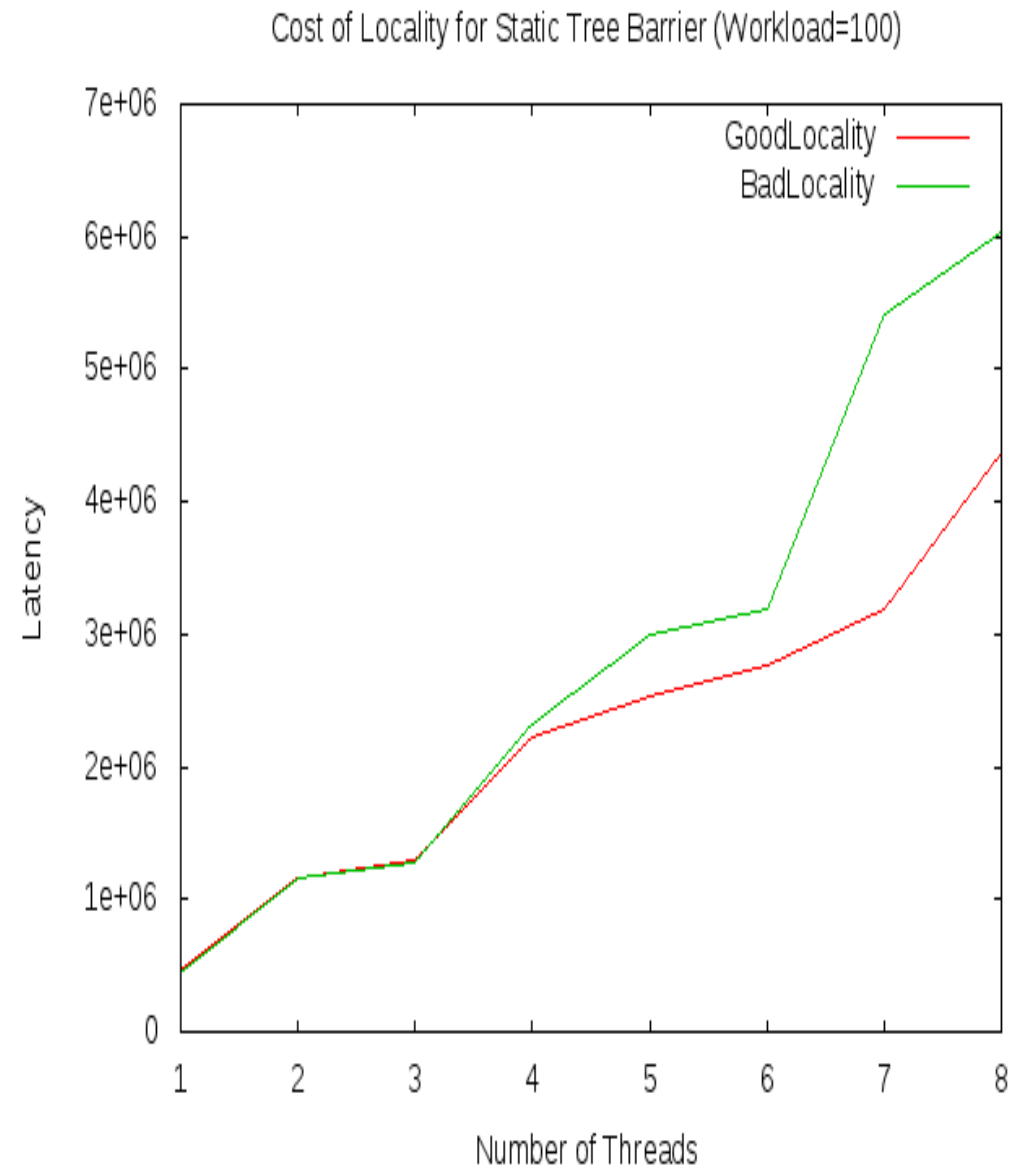
Experiments and Results (2/5)

- I first studied the effect of the relaxed memory model over the sequentially consistent one.
 - The cost of **mfence** for Intel
 - Relaxed is better!



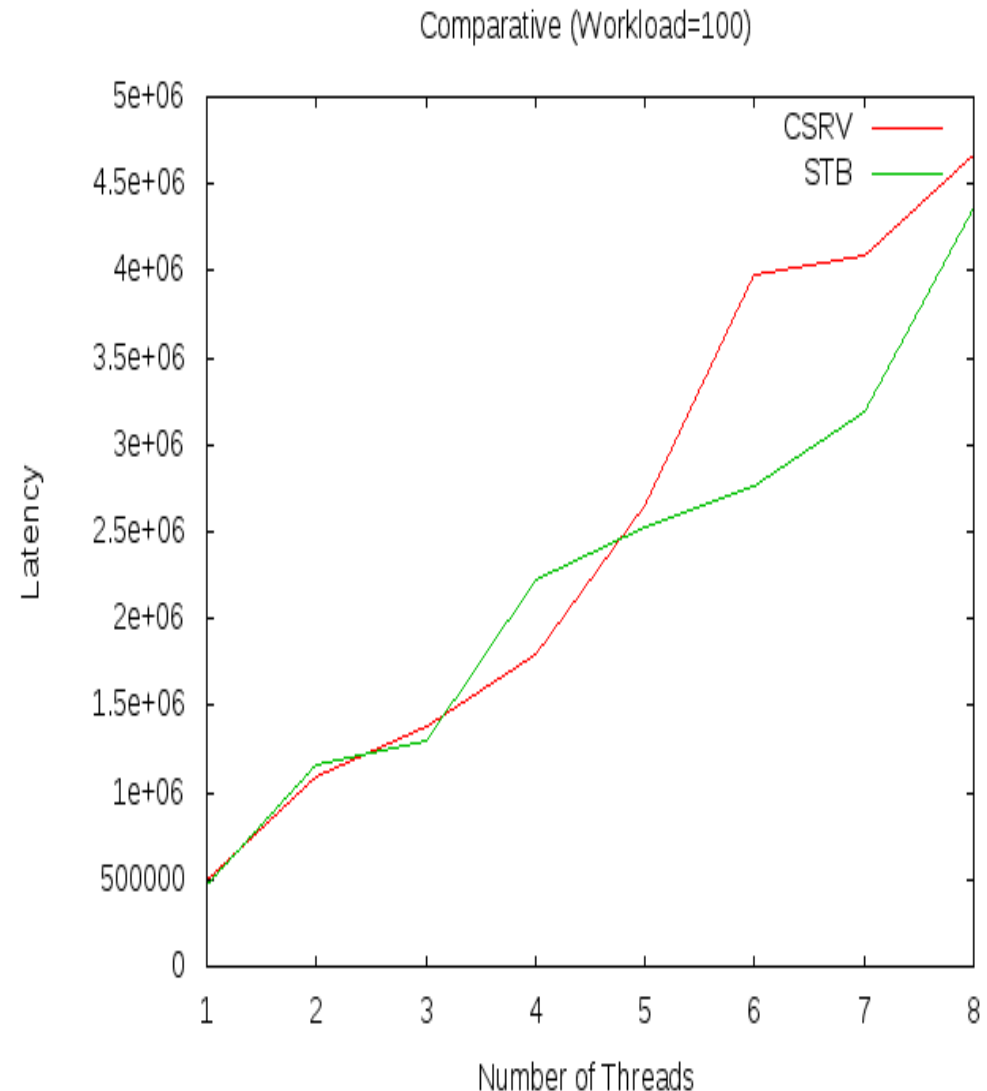
Experiments and Results (3/5)

- Then I studied the effect of *locality* for the static tree barrier
 - Locality matters!



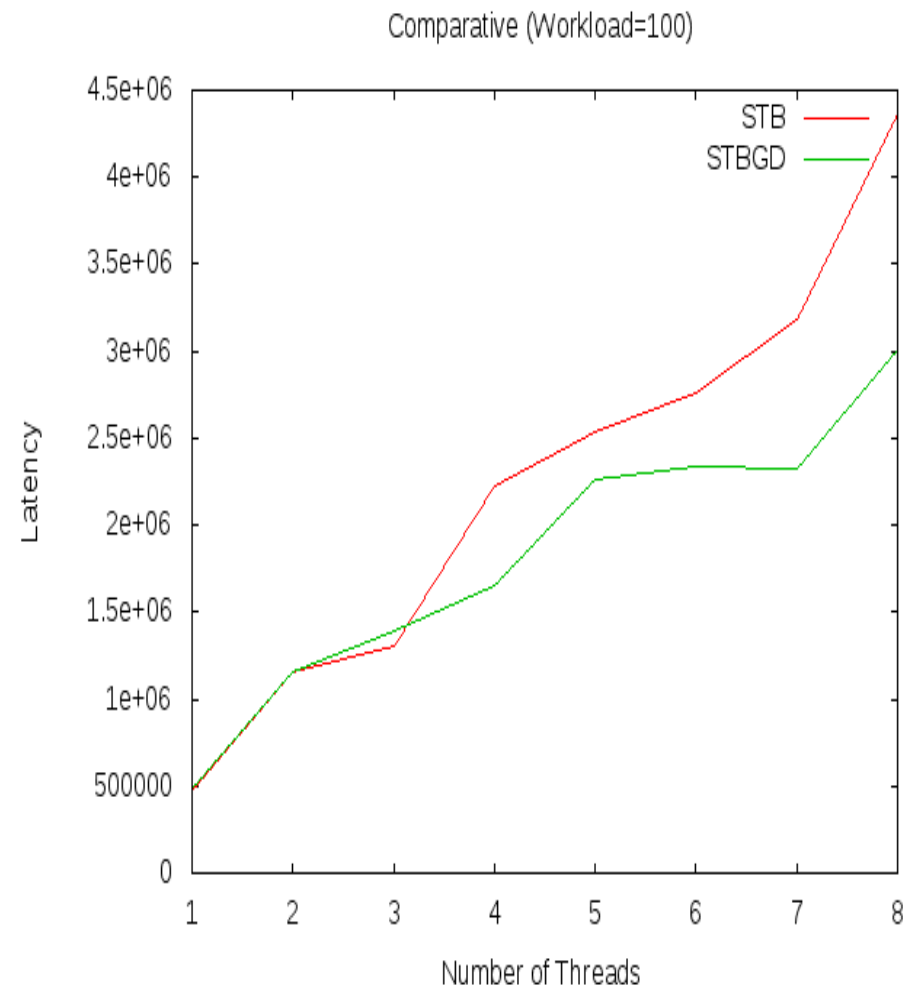
Experiments and Results (4/5)

- How do the Centralized Sense Reversing Barrier and Static Tree Barrier compare against?
 - Best versions
 - Relaxed+good-locality
 - Static Tree Barrier is better!



Experiments and Results (5/5)

- Surely the static tree barrier benefits from the *parallel* arrival phase versus the sequential bottleneck of the *shared counter* in the centralized sense reversing tree.
- But what about the departure phase?
 - I compare the previous static tree barrier with a static tree barrier whose global departure phase is implemented using a global sense flag
 - The version with the global departure flag is better.
 - The *ring interconnect* has to do something with it?
 - Further improvement over the centralized sense reversing barrier!



Conclusions

- OMITTED